

Levi Q Le

Tampa, FL | 305-778-2834 | qle1095@gmail.com | linkedin.com/in/levi-q-le

PROFESSIONAL SUMMARY

Staff Software Engineer and tech lead of a 12-engineer cross-functional team; led Data Fabric, the data platform central to Raft's \$450M acquisition by Leonardo DRS (announced July 2026). 6+ years building cloud-native infrastructure and data systems for government and financial-services customers, with deep expertise in distributed systems — Kubernetes, Kafka, Trino, Delta Lake — and full-stack product development in Go, Java, and TypeScript. Currently leading AI/LLM integration built entirely on self-hosted models — privacy-preserving multi-agent workflows where sensitive data never leaves the environment.

PROFESSIONAL EXPERIENCE

RAFT — Remote (Tampa, FL)

April 2022 – Present

Staff Software Engineer (2023 – Present)

Associate Software Engineer (2022 – 2023)

- Helped win a multi-million-dollar contract by building a customer-requested feature live during an executive demo.
- Go-to engineer for high-security delivery: lead deployments to air-gapped customer networks — rescuing several ahead of scheduled demos — and troubleshoot production systems on-site in restricted facilities, where debugging happens without internet access or standard tooling.
- Serve as the engineering counterpart to the customer, working with customer stakeholders to turn product requirements into technical designs.
- Mentor every engineer joining the team and interview candidates for all roles.
- Joined as employee #47 (now 350+); promoted from Associate directly to Staff in one year — skipping the Senior level — with multiple spot bonuses.
- Selected repeatedly for the hackathon teams behind Raft's competitive contract wins — programs renewed annually — including the one that became CBC2 (below).

Vulnerability Assessment Platform (Current)

- Own full-stack development of the platform security engineers use to assess release readiness: teams define a release from the Docker images it ships, and the Go/MongoDB backend pulls and aggregates CVE findings for each image automatically, replacing hand-compiled scanner reports.
- Built the TypeScript/React frontend — dashboards, drill-down views, and CVE justification workflows — giving analysts an auditable accept/reject record for each release decision.
- Automated risk scoring and surfaced critical vulnerabilities up front, enabling five analysts to review releases carrying tens of thousands of findings before they ship.
- Lead the AI integration, cutting the ~5-minute-per-CVE research burden: designed a multi-agent workflow on fully self-hosted models (Ollama, Gemma) — no customer data leaves their infrastructure — in which one agent drafts justifications from prior analyst-approved decisions and another curates them into a generalized knowledge base; every draft approved by a human analyst.

Data Fabric Platform

- Led the platform end to end — data engineering, software, and DevSecOps — onboarding ~10 data vendors and deploying to air-gapped customer networks.
- The platform anchors the multi-domain data-fusion capability cited in Leonardo DRS's \$450M acquisition of Raft (announced July 2026).
- Data Fabric is "Google over your data": users search thousands of datasets across 200+ sources, stream new data in with one click where supported, and see only the rows and columns their access level permits — serving hundreds of users.
- Built the catalog from scratch: ingestion via Kubernetes Operators/Controllers and CRDs (Kubebuilder), plus its web app — a Next.js/TypeScript frontend on a Go REST service with Cypress end-to-end tests.
- Integrated Kafka (via Strimzi), JupyterHub, MinIO, Delta Lake, Keycloak, Superset, OPA, Trino, and cert-manager into a single self-service data exploration stack.
- Implemented Attribute-Based Access Control (ABAC) for row- and column-level data filtering with custom-built Trino and Superset plugins and OPA rules.
- Wrote a dynamic Avro Kafka producer that infers schemas from upstream sources at consume time, letting users run SQL over Kafka data without per-topic configuration or manual schema maintenance — and authored the internal engineering blog post on the technique.

Pipeline & Workflow Automation

- Built pipeline jobs that automated previously manual deployment steps in the organization-wide CI/CD flow used by 30+ applications.

- Deployed Backstage with custom plugins that onboarded ~30 teams into the pipelines, giving each a single view of project status, security posture, and compliance readiness.
- Led pipeline observability: integrated GitLab APIs, S3, and Superset to publish metrics on jobs, security, and operations.

Communication Broker Channel Platform (CBC2)

- Designed Raft's initial architecture and built the technical foundation for CBC2 — a real-time data pipeline supporting mission-critical decision-making in the field — at contract kickoff, while continuing to lead Data Fabric.
- Designed and implemented the pipeline's ingest, validation, and enrichment stages with Kafka Streams in Scala, serving live data to consumers over WebSockets.
- Met the contract's throughput requirement of 10,000 messages per second by optimizing every stage of the pipeline and redesigning the components that became bottlenecks.

JPMORGAN CHASE — Tampa, FL

June 2018 – March 2022

Software Engineer (2020 – 2022)

Software Engineering Intern (2018, 2019)

- **Software Engineer Program (SEP):** Graduated the bank's two-year program in 18 months — fast-tracked — and mentored four junior engineers.
- **Service Container Platform:** Built four utilities (Messaging Toolkit, Messaging Failure Handler, Message Retry, Service Validation Suite) for a standardized platform whose shared library lets payment-processing teams adopt new tools without from-scratch integrations, and containerized legacy Java applications for the bank's cloud migration.
- **Billing & Pricing Platform:** Developed and maintained backend REST services for client charge processing, load-tested with JMeter and BlazeMeter.
- **Integrated Receivables (2019 internship):** Built a full-stack web application that onboards new clients to the program and version-tracks each configuration file.
- **Image-to-Text Ground Truth Collector (2018 internship):** Built a web app that collects human-labeled training data for an OCR model.

SKILLS & TECHNOLOGIES

Languages: Go, Java, Python, Scala, JavaScript/TypeScript, SQL, Bash

AI: LLM integration and tool calling (Ollama, Gemma, OpenAI API), multi-agent workflows

Platform & Cloud: Docker, Kubernetes (Operators, CRDs, Kubebuilder), Helm, Istio, AWS (EKS, S3), Azure (AKS), MinIO, Linux

Architecture: Distributed systems, microservices, event-driven streaming, ETL

Data & Streaming: Kafka (Kafka Streams, Kafka Connect, Strimzi), Avro, Trino, Delta Lake, Superset, PostgreSQL, MongoDB

Backend & Frontend: Spring Boot, Quarkus, REST APIs, WebSockets, React, Next.js, Material UI

Security & Identity: Keycloak, OPA/Rego, OAuth, SSO, ABAC, DevSecOps

Testing & CI/CD: JUnit, Mockito, Hamcrest, Cypress, JMeter, GitLab CI/CD, Backstage, Git

Spoken Languages: English (bilingual proficiency), Vietnamese (native)

CERTIFICATIONS & CLEARANCES

- Certified Kubernetes Application Developer (CKAD), 2025
- CompTIA Security+ (SY0-701)
- CompTIA SecurityX (CAS-005) — in progress
- Active U.S. Government TS/SCI Clearance (2022 – Present)

EDUCATION

Bachelor of Science, IT Software Development — Florida International University, Miami, FL (2019) — *magna cum laude*

SIDE PROJECTS

Tea Shop Operations Automation Platform

- Built a full-stack automation system for a local tea shop — in production, handling up to 100 orders/day — linking Square POS orders to standardized recipes and eliminating recipe recall errors during preparation.
- Replaced verbal order-taking with a self-service kiosk and real-time preparation queue, and built employee procedure apps enforcing step-by-step recipes — cutting new-hire ramp-up from about a month to a few days and reducing operating costs.

AI Health Coach — Multi-Agent Fitbit Assistant

- Built with the Google Agent Development Kit and OpenAI API: nutrition, exercise, and sleep agents collaborate on Fitbit data; users chat with each expert to log meals and workouts for personalized coaching.